

Cheng Chang

chc012@alumni.stanford.edu | chc012.github.io | linkedin.com/in/chc012 | github.com/chc012

Education

Stanford University <i>MS in Computational and Mathematical Engineering</i>	Stanford, CA, USA Sep 2022 – Jun 2024
---	--

University of California San Diego <i>BS in Cognitive Science; BA in Philosophy</i>	La Jolla, CA, USA Sep 2018 – Jun 2022
---	--

Experience

Uniphore (acquired Orby AI) <i>Senior AI Scientist</i>	Palo Alto, CA, USA Feb 2026 – Present
--	--

- **AI Research:** Drive research on software-engineering agent evaluation and LLM-as-a-judge meta-evaluation.
- **Software Engineering:** Lead development of a production-grade system that evaluates AI agents end-to-end using an LLM-as-a-judge approach.

<i>Senior Software Engineer</i>	Aug 2025 – Jan 2026
---------------------------------	---------------------

- **AI Research:** Developed coding agents for search automation; created an open-source GUI automation benchmark and trained web navigation models.
- **Software Engineering:** Built services to help users select small language models (SLMs) for custom enterprise use cases.

Orby AI <i>Software Engineer</i>	Mountain View, CA, USA Jul 2024 – Aug 2025
--	---

- **Multimodal Model Research:** Trained generative models and built web automation agents achieving state-of-the-art results on public benchmarks.
- **Software Engineering:** Integrated LLMs into backend automation tools for code generation and web interaction.

<i>Software Engineer Intern</i>	Apr 2024 – Jun 2024
---------------------------------	---------------------

- Built a distributed web UI data crawler for foundation model training.

SALT Lab, Stanford University <i>Graduate Student Researcher</i>	Stanford, CA, USA Oct 2023 – Mar 2024
--	--

- Developed CARE 2.0 (LLM-based counselor training platform); implemented UI, backend APIs, and database.

GrubMarket, Inc. <i>Data Science Intern</i>	San Francisco, CA, USA Jun 2023 – Aug 2023
---	---

- Designed and implemented recommender systems for fresh produce.

Selected Publications

PADMÉ: Preference Alignment Data Synthesis for Meta-Evaluation of LM Agent Evaluators. *NeurIPS TAE Workshop 2026 (to appear).*

WARC-Bench: Web Archive Based Benchmark for GUI Subtask Executions. *ICLR 2026.*

[WebDreameer] Is Your LLM Secretly a World Model of the Internet? Model-Based Planning for Web Agents. *TMLR 2025.*

[UGround] Navigating the Digital World as Humans Do: Universal Visual Grounding for GUI Agents. *ICLR 2025.*

Roleplay-doh: Enabling Domain Experts to Create LLM-Simulated Patients. *EMNLP 2024.*

Tech Stacks

Languages: Python, C/C++, TypeScript/JavaScript, R, SQL, HTML, CSS, TeX/LaTeX

ML & LLM: PyTorch, Hugging Face, LiteLLM, OpenAI, Anthropic, Fireworks AI

Infrastructure & Tools: AWS, Kubernetes/Docker, Ray, Databricks, Linux, Git, FastAPI, Flask